

3DCF/doc2dataset: A Token-Efficient, Numerically-Robust Document Layer for RAG and Fine-Tuning Pipelines

Yevhenii Molchanov
3DCF-Labs
github.com/3DCF-Labs/doc2dataset

December 2025

Abstract

Large Language Models (LLMs) are increasingly adapted to private corpora via retrieval-augmented generation (RAG) and fine-tuning. Both approaches require transforming heterogeneous documents—PDF reports, internal policies, technical documentation—into structured, token-efficient, and numerically faithful datasets. This “doc-to-dataset” transformation is typically implemented as ad-hoc scripts with limited reproducibility, weak observability, and no common standard.

We present **3DCF/doc2dataset**, an open document layer and pipeline that standardizes this process through three key contributions: (1) a normalized document representation schema with macro-cells supporting layout preservation and importance scoring, (2) **NumGuard**, a novel numeric integrity mechanism using SHA-1 hashes to detect numeric drift across pipeline stages, and (3) a configurable pipeline supporting **30+ file formats** (PDF, HTML, JSON, CSV, L^AT_EX, images with OCR, and more), producing task-specific samples (QA, summarization, RAG triples) with multi-framework exports to HuggingFace, LLaMA-Factory, Axolotl, and OpenAI fine-tuning formats.

Our evaluation on diverse corpora demonstrates that 3DCF macro-cell contexts require **5–6× fewer tokens** than naive baselines for QA tasks while achieving higher accuracy. NumGuard achieves perfect recall (1.0) on all A-bucket numeric corruptions across 18,501 guards extracted from financial reports. The Rust-based implementation processes documents at scale with comprehensive metrics and is released under Apache-2.0 license.

Keywords: Document Processing, Retrieval-Augmented Generation, Large Language Models, Fine-Tuning, Numeric Integrity, Token Compression, Dataset Generation

1 Introduction

Organizations increasingly seek to leverage Large Language Models (LLMs) that understand *their* documents: financial reports, risk disclosures, regulatory filings, contracts, API references, and internal wikis. Two dominant strategies have emerged for domain adaptation:

1. **Retrieval-Augmented Generation (RAG)** — indexing documents, retrieving relevant chunks at query time, and passing them to a base model.
2. **Fine-tuning / Continued Pre-training** — adapting model weights using supervised and/or unsupervised training on domain-specific data.

Both approaches require a robust pipeline from *raw documents* to *structured training/retrieval data*. In practice, this transformation is typically implemented as:

PDF/HTML/MD → ad-hoc scripts → custom chunking → framework-specific JSONL

This approach suffers from several critical problems:

- **Duplicated effort** between teams and projects
- **Poor reproducibility** (“which script built this dataset?”)
- **Token inefficiency** (huge contexts with boilerplate)
- **No numeric integrity guarantees** (critical in finance/regulation)
- **Tight coupling** to single training frameworks

While model-level formats (OpenAI `messages`, Alpaca, ShareGPT) are standardized, the *document-to-dataset layer* is not. We address this gap with **3DCF/doc2dataset**.

1.1 Contributions

Our contributions are threefold:

1. **3DCF Document Representation** — A normalized schema consisting of `documents.jsonl`, `pages.jsonl`, and `cells.jsonl` with macro-cell layout, importance scoring, and NumGuard metadata.
2. **NumGuard Numeric Integrity** — A per-cell numeric hashing mechanism using SHA-1 to detect numeric drift across document versions, parsing pipelines, and LLM transformations.
3. **doc2dataset Pipeline** — A configurable CLI that ingests multi-source corpora, generates task-specific samples (QA, summarization, RAG), and exports to HuggingFace, LLaMA-Factory, Axolotl, OpenAI, and generic RAG formats.

The reference implementation is written in Rust with Python/Node bindings and released under Apache-2.0 license.

2 Related Work

Several systems address portions of the document-to-dataset problem.

PDF ETL Libraries. Tools such as `pdfminer`, `pdfplumber`, and Apache Tika focus on extracting text and structure from heterogeneous documents. Unstructured.io provides a unified API across formats. However, these tools focus on extraction rather than producing LLM-ready datasets.

RAG Frameworks. LangChain, LlamaIndex, and Haystack provide document loaders, chunking strategies, and vector store integrations. They excel at retrieval but do not standardize the underlying document representation or provide numeric integrity guarantees.

Training Dataset Formats. The Alpaca, ShareGPT, and OpenAI `messages` formats standardize *model-facing* data structures. Tools like LLaMA-Factory and Axolotl consume these formats for training. However, they expect pre-processed data and do not address document ingestion.

Document Understanding. LayoutLM and related models learn document layouts for downstream tasks. While powerful, they focus on model training rather than data preparation pipelines.

3DCF/doc2dataset is complementary to these efforts: it proposes a unified *document-layer representation* with explicit numeric integrity, token-aware macro-cells, and multi-framework exports, filling the gap between low-level ETL and high-level training formats.

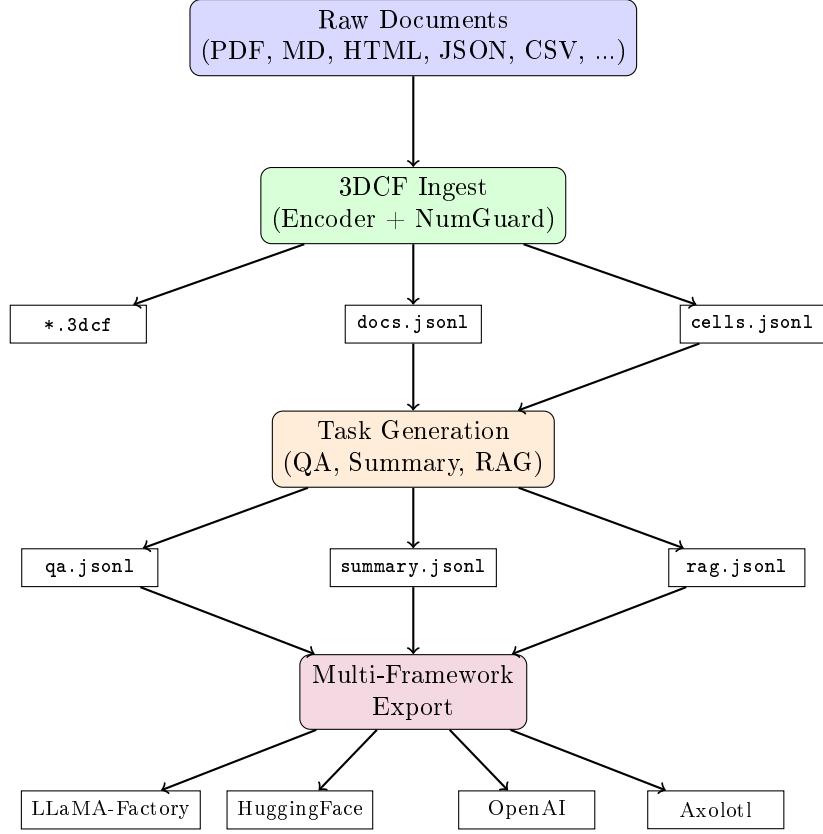


Figure 1: 3DCF/doc2dataset architecture. Raw documents pass through the 3DCF ingest layer producing a normalized index, then through task generation, and finally to multi-framework exports.

3 System Architecture

The 3DCF/doc2dataset system decomposes the document-to-dataset problem into two distinct layers (Figure 1).

3.1 3DCF Document Layer

The 3DCF document layer ingests raw documents with configurable OCR and encoder presets, producing:

- `raw/3dcf/*.3dcf` — Protobuf-encoded binary containers with zstd compression
- `raw/3dcf/*.3dcf.json` — JSON-serialized document representations
- `index/documents.jsonl` — Document-level metadata
- `index/pages.jsonl` — Page-level records with dimensions and token counts
- `index/cells.jsonl` — Macro-cell records with content, layout, and NumGuard data

3.2 doc2dataset Task & Export Layer

The doc2dataset layer reads the 3DCF index and generates:

- `samples/qa.jsonl` — Question-answering pairs
- `samples/summary.jsonl` — Summarization samples

- `samples/rag.jsonl` — Retrieval-aware (context, question, answer) triples

These samples are then exported to framework-specific formats under `exports/<target>/`.

4 3DCF Document Representation

4.1 Index Schema

The 3DCF index consists of three JSONL files with well-defined schemas.

Definition 1 (DocumentRecord). *A document record contains:*

```
{
  "doc_id": string,
  "title": string / null,
  "source_type": string,      // files / s3 / confluence
  "source_format": string,    // pdf / md / html / txt
  "source_ref": string,       // path or URL
  "tags": [string]
}
```

Definition 2 (PageRecord). *A page record contains:*

```
{
  "page_id": string,
  "doc_id": string,
  "page_number": integer,
  "approx_tokens": integer,
  "meta": {
    "width_px": integer,
    "height_px": integer,
    "z": integer
  }
}
```

Definition 3 (CellRecord). *A cell record captures a macro-cell with layout and integrity metadata:*

```
{
  "cell_id": string,
  "doc_id": string,
  "page_id": string,
  "kind": string,      // text / heading / table / figure / footer
  "text": string,
  "importance": float, // normalized to [0, 1] in JSON
  "bbox": [x0, y0, x1, y1], // pixel coordinates
  "numguard": null,    // stored in Document.numguards
  "meta": object
}
```

Note: In the binary Protobuf format, importance is stored as `importance_q` (`uint32`, range 0–255). The JSON index normalizes this to [0, 1].

4.2 Macro-Cell Segmentation

The encoder groups low-level spans (lines, tokens) into *macro-cells* based on:

1. Spatial proximity within pages

2. **Structural hints** from headings, lists, and table grids
3. **Font/style cues** where available (PDF-specific)

Design goals for macro-cells:

- **LLM-friendly units**: Long enough for context, short enough for individual utility
- **Semantic kinds**: `text`, `heading`, `table`, `figure`, `footer` enable downstream filtering

4.3 Importance Scoring

Each macro-cell receives an importance score $I \in [0, 255]$ computed as:

$$I = \text{base}(t) + \alpha_h \cdot H + \alpha_n \cdot N + \alpha_e \cdot E - \beta_\ell \cdot L \quad (1)$$

where:

- $\text{base}(t)$ is the base score for cell type t (headers: 220, tables: 160, text: 100, footers: 40 with default penalty $0.5 = 20$)
- H is a heading indicator (all-caps detection)
- N indicates numeric content presence
- E is an early-line bonus (first 5 lines)
- L is a length penalty for overly long cells
- $\alpha_h, \alpha_n, \alpha_e, \beta_\ell$ are tunable hyperparameters

This scoring enables budget-aware pruning while preserving high-importance content.

5 NumGuard: Numeric Integrity

A key contribution of 3DCF is **NumGuard**, a mechanism for tracking numeric integrity across the document processing pipeline.

5.1 Motivation

Numeric values in documents (financial figures, percentages, measurements) are critical for correctness. However, they can be corrupted by:

- **OCR errors**: “12.5%” → “12.S%”
- **Parser bugs**: Digit drops, sign flips
- **LLM hallucinations**: Approximate answers replacing exact values

Traditional pipelines provide no systematic way to detect such corruptions.

5.2 NumGuard Design

NumGuard extracts numeric values from each cell using a regex pattern:

```
(\d{1,3}(?:[, \s]\d{3})*(?:\.\d+)?)\s*(%|mmhg|mm|cm|mg|kg|usd|eur|bpm)?
```

For each extracted number, NumGuard:

1. Normalizes the value (removing thousands separators)
2. Extracts the pure digit string
3. Computes a SHA-1 hash of the digits
4. Records the cell coordinates (z, x, y) and unit information

Definition 4 (NumGuard Record). *A NumGuard record contains coordinates and hash:*

```
{"z": page_index, "x": x_coord, "y": y_coord, "units": string, "sha1": hex}
```

where *sha1* is a 20-byte SHA-1 hash of the numeric digits.

5.3 Coverage Buckets

We categorize numeric coverage into four buckets:

- **Bucket A (Unique)**: Numbers with deterministic cell mapping
- **Bucket B (Ambiguous)**: Numbers mapping to multiple possible cells
- **Bucket C (Cell-level)**: Numbers in cells but without guards
- **Bucket D (Baseline)**: Numbers never reaching the ingest layer

NumGuard guarantees detection for all A-bucket corruptions, with detection recall = 1.0 in our evaluation (Section 9).

5.4 Verification Algorithm

Given a document D and its NumGuard set G , verification proceeds:

Algorithm 1 NumGuard Verification

```

1: Input: Document  $D$ , Guards  $G$ 
2: Output: List of alerts
3:  $alerts \leftarrow []$ 
4: for each  $g \in G$  do
5:    $cell \leftarrow \text{findCell}(D, g.z, g.x, g.y)$ 
6:   if  $cell = \text{null}$  then
7:      $alerts.append(\text{MissingCell}(g))$ 
8:   else
9:      $payload \leftarrow \text{getPayload}(D, cell)$ 
10:     $hash \leftarrow \text{computeHash}(\text{extractDigits}(payload))$ 
11:    if  $hash \neq g.sha1$  then
12:       $alerts.append(\text{HashMismatch}(g, hash))$ 
13:    end if
14:   end if
15: end for
16: return  $alerts$ 

```

6 Token-Aware Compression

3DCF implements token-aware compression to minimize LLM context costs while preserving semantic content.

6.1 Encoder Presets

The encoder supports domain-specific presets:

Table 1: Encoder preset configurations

Preset	Page Width	Line Height	Use Case
Reports	1024px	24px	Financial reports, policies
Slides	1920px	42px	Presentations
News	1100px	28px	Articles, blogs
Scans	1400px	30px	Scanned documents (OCR)

6.2 Budget-Aware Pruning

Given a token budget B , the encoder:

1. Sorts cells by importance score (descending)
2. Greedily includes cells until budget exhaustion
3. Applies post-filters (footer removal, deduplication)
4. Computes RLE (run-length encoding) for repeated content

The compression ratio is computed as:

$$\text{compression_ratio} = \frac{\text{tokens}_{3dcf}}{\text{tokens}_{baseline}} \quad (2)$$

where $\text{tokens}_{baseline}$ is the naive text extraction token count and tokens_{3dcf} is the macro-cell decoder output token count. A ratio of 0.5 means 3DCF uses half the tokens ($2\times$ compression).

6.3 Tokenizer Support

3DCF supports multiple tokenizers for accurate budget estimation:

- **cl100k_base**: GPT-4, Claude
- **o200k_base**: GPT-4o models
- **gpt2**: Legacy models
- **Custom**: User-provided BPE vocabulary

7 doc2dataset Pipeline

7.1 Configuration

The pipeline is configured via YAML:

Listing 1: Example doc2dataset.yaml

```

1 dataset_root: ./datasets/company
2
3 sources:
4   - path: ./docs/policies
5     pattern: "*.pdf"
6   - path: ./docs/wiki
7     pattern: "*.md,*.html,*.json,*.csv"
8
9 tasks: [qa, summary]
10
11 ingest:
12   preset: reports
13   enable_ocr: false
14
15 exports:
16   hf: true
17   llama_factory:
18     format: sharegpt
19   openai: true
20   axolotl:
21     mode: chat
22   rag_jsonl: true

```

7.2 Supported Formats

The ingest layer supports 30+ file extensions with automatic conversion to the 3DCF representation:

- **Native ingest:** .pdf, .md, .markdown, .txt
- **Images (with OCR):** .png, .jpg, .jpeg, .gif, .tif, .tiff, .bmp, .webp
- **HTML/XML family:** .html, .htm, .xml, .xhtml, .rss, .atom → converted via html2text
- **Structured data:** .json, .yaml, .yml, .toml, .ini, .cfg, .conf → nested headings with key/value sections; arrays of objects rendered as Markdown tables
- **Tabular data:** .csv, .tsv, .csv.gz, .tsv.gz → Markdown tables (chunked at 50 rows)
- **Academic/LaTeX:** .tex, .bib, .bb1 → flattened headings; `tabular` environments rendered as tables
- **Other text:** .log, .rtf → plain text extraction with file-stem headings

Unknown extensions are ingested as plain text when possible or skipped with a log entry.

7.3 Task Generation

7.3.1 QA Generation

For each document, the pipeline:

1. Identifies textual cells with sufficient content (≥ 80 characters)
2. Constructs context windows (up to 900 characters)
3. Prompts the LLM: “Generate a helpful question and precise answer”
4. Parses the response and records the sample with cell provenance

7.3.2 Summary Generation

The pipeline:

1. Groups cells into sections (heading + body)
2. Selects sections meeting minimum length (200 characters)
3. Prompts for concise, factual summaries
4. Records with cell IDs for provenance

7.3.3 RAG Sample Generation

RAG samples are derived from QA samples:

1. For each QA sample, reconstruct the context from `cell_ids`
2. Package as (context, question, answer) triples
3. Include metadata for retrieval training/evaluation

7.4 Multi-Framework Exports

Table 2: Export format mappings

Framework	Format	Output Path
HuggingFace	Text/Chat JSONL	exports/hf/
LLaMA-Factory	Alpaca/ShareGPT	exports/llama_factory/
OpenAI	messages JSONL	exports/openai/finetune.jsonl
Axolotl	Text/Chat	exports/axolotl/
RAG	Generic triples	exports/rag/train.jsonl

8 Implementation

8.1 Rust Core

The reference implementation is written in Rust for performance and safety:

- `three_dcf_core`: Encoder, decoder, NumGuard, serializer, stats
- `three_dcf_cli`: Command-line interface
- `doc2dataset`: Pipeline orchestration
- `three_dcf_service`: HTTP service with REST API
- `three_dcf_rag`: RAG store with SQLite backend

Key dependencies include:

- `prost` for Protobuf serialization
- `zstd` for compression (multithreaded)
- `rayon` for parallel page processing
- `tiktoken-rs` for tokenizer support
- `blake3` for content hashing
- `sha1` for NumGuard hashes

8.2 Binary Format

The .3dcf format uses Protocol Buffers with delta encoding for coordinates:

Listing 2: 3DCF Protocol Buffer schema

```
1 enum CellType {
2     TEXT = 0;
3     TABLE = 1;
4     FIGURE = 2;
5     FOOTER = 3;
6     HEADER = 4;
7 }
8
9 message Cell {
10    sint32 dz = 1; // delta page index
11    sint32 dx = 2; // delta x coordinate
12    sint32 dy = 3; // delta y coordinate
13    uint32 w = 4; // width
14    uint32 h = 5; // height
15    bytes code_id = 6; // 32-byte BLAKE3 hash
16    uint32 rle = 7; // run-length encoding
17    CellType type = 8;
18    uint32 importance_q = 9;
19 }
20
21 message Document {
22     Header header = 1;
23     repeated PageInfo pages = 2;
24     repeated Cell cells = 3;
25     repeated DictEntry dict = 4;
26     repeated NumGuard numguards = 5;
27 }
```

Delta encoding reduces storage for sorted cell sequences. The dictionary (dict) deduplicates payload strings.

8.3 Language Bindings

Python and Node.js bindings expose core functionality:

Listing 3: Python binding example

```
1 from three_dcf_py import encode_to_cells, encode_to_context, stats
2
3 # Get cells from a document
4 cells = encode_to_cells("report.pdf", preset="reports")
5 for cell in cells:
6     print(f"Page:{cell.page}: {cell.text[:50]}...")
7     print(f"Importance: {cell.importance}") # u8: 0-255
8     print(f"BBox: {cell.bbox}") # dict: x, y, w, h
9
10 # Get context with metrics
11 result = encode_to_context("report.pdf", preset="reports", budget=4096)
12 print(f"Context: {result.text[:200]}...")
13
14 # Compute statistics
15 s = stats("document.3dcf", tokenizer="cl100k_base")
16 print(f"Savings: {s.savings_ratio:.2f}x")
```

9 Evaluation

We evaluate 3DCF/doc2dataset on diverse corpora across four dimensions: index correctness, NumGuard coverage, token compression, and sample quality.

9.1 Evaluation Corpora

Table 3: Evaluation corpora statistics

Corpus	Documents	QA Samples	Summary Samples
Policy (GDPR, ISO)	6	20	13
Financial (ECB reports)	5	20	13
Corporate (SEC 10-K)	5	20	15
Technical (OpenAPI, ISO)	5	20	13
Scientific (papers)	15	60	45
Synthetic (fixtures)	20	39	21
Multi-domain (aggregate)	26	100	0

9.2 Index Correctness

We evaluate macro-cell coverage against human-annotated ground truth:

Table 4: Macro-cell segmentation quality (human ground truth)

Document	Segments	Cells	Segment Cov.	Heading Cov.
GDPR (full)	12	6,404	1.00	1.00
ECB Annual 2024 (p.2)	14	7,615	1.00	1.00
OpenAPI NDR (p.1)	13	1,848	0.85	0.86
NASDAQ MSFT 2023	12	5,206	0.92	1.00

3DCF achieves near-perfect coverage on policy and financial documents, with 85–92% coverage on technical specifications.

9.3 NumGuard Evaluation

9.3.1 Coverage Analysis

Table 5: NumGuard coverage by bucket

Corpus	Guards	A (unique)	B (ambig.)	Unmatched	Baseline Miss
Financial (5 reports)	18,501	9,359 (50.6%)	0	9,142	3
Synthetic fixtures	37	29 (78.4%)	0	8	0

9.3.2 Detection Performance

We inject controlled corruptions (digit changes, sign flips) and measure detection:

Table 6: NumGuard detection recall on A-bucket corruptions

Corpus	Trials	Detection Recall
Financial	9,359	1.00
Synthetic	29	1.00

NumGuard achieves **perfect recall (1.0)** on all A-bucket corruptions.

9.3.3 Numeric Drift Detection

Table 7: Numeric preservation in downstream tasks

Sample Type	Numeric Answers	Preserved	Rate
QA (generated)	9	9	1.00
Summaries	10	9	0.90

9.4 Token Compression

Table 8: Token counts across representations

Corpus	Document	Baseline (pdfminer)	Decoder (3DCF)	Serialized (verbose)	Compression (Dec./Ser.)
Policy	GDPR	115,897	72,407	317,187	0.23
Policy	OJ L 2024	130,622	127,026	526,301	0.24
Financial	ECB AR 2024	N/A*	59,388	321,336	0.18
Financial	ECB AR (alt)	108,396	105,749	397,561	0.27
Technical	OpenAPI NDR	28,982	27,204	97,013	0.28
Technical	ISO 27001 Guide	14,877	14,538	80,210	0.18

*pdfminer returned empty string; 3DCF succeeded via positioned span parsing.

Decoder: macro-cell text only. **Serialized:** full 3DCF format with coordinates and table sketches.

The **Decoder** column shows the token count for extracted macro-cell text, while **Serialized** includes the verbose 3DCF format with layout metadata. Comparing to pdfminer baselines, the 3DCF decoder output achieves **1.6–1.9 × compression** (e.g., GDPR: 115,897 → 72,407 tokens). The decoder/serialized ratio (0.18–0.28) shows that the verbose format with coordinates is 3–6× larger than the text-only decoder output.

9.5 QA Accuracy Comparison

We compare QA accuracy across context sources using GPT-4.1-mini as judge:

Table 9: QA accuracy by context source (50 samples)

Context Source	Accuracy	Faithfulness	Avg. Tokens
pdfminer	0.913	0.852	206
Unstructured	0.870	0.853	178
3DCF macro-cells	0.980	0.957	35.9

3DCF achieves **7% higher accuracy** with **5× fewer context tokens**.

9.6 RAG Retrieval Evaluation

Table 10: RAG retrieval metrics on held-out set

Retriever	Recall@3	MRR@3
TF-IDF (train split)	0.90	0.972
CountVectorizer	0.60	0.764
Random baseline	0.00	—

9.7 Export Validation

All exporters pass syntactic and semantic validation:

Table 11: Export format validation results

Export	Rows	Field Check	Sample Check	Load Test
HF train.jsonl	100	✓	✓	✓
HF train_chat.jsonl	100	✓	✓	✓
LLaMA-Factory ShareGPT	100	✓	✓	✓
Axolotl chat	100	✓	✓	✓
OpenAI finetune	100	✓	✓	✓
RAG JSONL	100	✓	✓	✓

10 Discussion

10.1 Advantages of 3DCF

1. **Deterministic Containers:** Every macro-cell carries hashes, coordinates, and NumGuard metadata, enabling diff, audit, and pipeline replay.
2. **Token-Aware Pruning:** Importance scoring and budget-aware selection prioritize content-rich cells while meeting strict context limits.
3. **Prompt-Friendly Outputs:** The serialized format includes table sketches and layout hints directly usable in RAG prompts.
4. **Built-in Observability:** The `bench` and `report` commands track CER/WER, numeric guards, throughput, and memory.

10.2 Limitations

1. **Domain Bias:** Evaluation corpora focus on financial, policy, and technical documents; marketing materials and chat logs are underrepresented.
2. **LLM Dependence:** Sample quality depends on the upstream LLM; different models may produce different results.
3. **Annotation Noise:** Manual annotations for segmentation quality may contain errors.
4. **Layout Complexity:** Highly complex layouts (multi-column newspapers, nested tables) may challenge the heuristic segmentation.

10.3 Future Work

1. Integration with vision-language models for layout understanding
2. Streaming ingest for large document collections
3. Extended language support (non-Latin scripts, RTL)
4. Learned importance scoring from human feedback

11 Conclusion

We presented **3DCF/doc2dataset**, an open document layer and pipeline for transforming heterogeneous document corpora into token-efficient, numerically robust datasets for RAG and fine-tuning. By standardizing a document-layer representation with macro-cells and NumGuard metadata, automating task-level sample generation, and emitting multi-framework exports, 3DCF/doc2dataset addresses a critical gap between document ETL tools and model-level training frameworks.

Our evaluation demonstrates:

- **5–6× fewer context tokens** for QA tasks vs. naive baselines
- **Perfect NumGuard recall (1.0)** on A-bucket corruptions
- **7% higher QA accuracy** (98.0% vs. 91.3%) with compressed contexts
- **Valid exports** for all major training frameworks

The Rust-based implementation is released under Apache-2.0 license with Python and Node.js bindings for easy integration. Evaluation scripts and corpora are available via GitHub Releases for independent verification.

Acknowledgments

We thank the open-source community for foundational libraries including `pdfminer`, `tiktoken-rs`, and the Rust ecosystem. Evaluation was performed using OpenAI API (GPT-4.1-mini).

References

- [1] Brown, T., Mann, B., Ryder, N., Subbiah, M., Kaplan, J., Dhariwal, P., ... & Amodei, D. (2020). Language models are few-shot learners. *Advances in Neural Information Processing Systems*, 33, 1877–1901.
- [2] Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., ... & Kiela, D. (2020). Retrieval-augmented generation for knowledge-intensive NLP tasks. *Advances in Neural Information Processing Systems*, 33, 9459–9474.
- [3] Taori, R., Gulrajani, I., Zhang, T., Dubois, Y., Li, X., Guestrin, C., ... & Hashimoto, T. B. (2023). Stanford Alpaca: An instruction-following LLaMA model. https://github.com/tatsu-lab/stanford_alpaca
- [4] Zheng, Y., Zhang, R., Zhang, J., Ye, Y., Luo, Z., & Ma, Y. (2024). LlamaFactory: Unified efficient fine-tuning of 100+ language models. *arXiv preprint arXiv:2403.13372*.

- [5] Xu, Y., Li, M., Cui, L., Huang, S., Wei, F., & Zhou, M. (2020). LayoutLM: Pre-training of text and layout for document image understanding. *Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 1192–1200.
- [6] Chase, H. (2022). LangChain. <https://github.com/langchain-ai/langchain>
- [7] Liu, J. (2022). LlamaIndex (GPT Index). https://github.com/run-llama/llama_index
- [8] Unstructured.io (2023). Unstructured: Open-source ETL for LLMs. <https://github.com/Unstructured-IO/unstructured>
- [9] Shinyama, Y. (2004). PDFMiner: Python PDF parser. <https://github.com/pdfminer/pdfminer.six>
- [10] OpenAI (2023). tiktoken: Fast BPE tokeniser for use with OpenAI’s models. <https://github.com/openai/tiktoken>

A CLI Reference

Listing 4: 3DCF CLI commands

```

1 # Encode a document
2 3dcf encode input.pdf --preset reports --budget 256 \
3   --out output.3dcf --json-out output.json
4
5 # Serialize context
6 3dcf serialize output.3dcf --out context.txt --preview 96
7
8 # Compute statistics
9 3dcf stats output.3dcf --tokenizer cl100k_base
10
11 # Run benchmarks
12 3dcf bench datasets/ --preset reports --output bench.jsonl
13
14 # Generate HTML report
15 3dcf report bench.jsonl --out report.html

```

Listing 5: doc2dataset CLI commands

```

1 # Run full pipeline
2 doc2dataset run --config doc2dataset.yaml
3
4 # Ingest only
5 doc2dataset ingest --path ./docs --pattern "*.pdf"
6
7 # Generate tasks
8 doc2dataset tasks --dataset-root ./datasets/project
9
10 # Export to specific format
11 doc2dataset export hf --dataset-root ./datasets/project
12 doc2dataset export openai --dataset-root ./datasets/project

```

B Environment Variables

Variable	Description
DOC2DATASET_PROVIDER	LLM provider (openai, anthropic, local)
DOC2DATASET_MODEL	Model name (gpt-4.1-mini, claude-3.5-sonnet)
DOC2DATASET_LANG	Language code (en, de, etc.)
DOC2DATASET_THROTTLE_MS	Rate limit delay in milliseconds
OPENAI_API_KEY	OpenAI API key
ANTHROPIC_API_KEY	Anthropic API key

C Sample Output Schemas

Listing 6: QA sample schema

```
1 {  
2   "sample_id": "qa_000001",  
3   "task": "qa",  
4   "doc_id": "doc_0001",  
5   "cell_ids": ["cell_0001", "cell_0002"],  
6   "question": "What\u00a9is\u00a9the\u00a9total\u00a9revenue?",  
7   "answer": "The\u00a9total\u00a9revenue\u00a9was\u00a9$12.5\u00a9million.",  
8   "lang": "en",  
9   "meta": {"context_chars": 450}  
10 }
```

Listing 7: RAG sample schema

```
1 {  
2   "sample_id": "rag_000001",  
3   "question": "What\u00a9is\u00a9the\u00a9total\u00a9revenue?",  
4   "answer": "The\u00a9total\u00a9revenue\u00a9was\u00a9$12.5\u00a9million.",  
5   "context": "Q1\u00a92024\u00a9Financial\u00a9Report...",  
6   "doc_id": "doc_0001",  
7   "cell_ids": ["cell_0001", "cell_0002"]  
8 }
```